

FlashDeflate: Achieving $2\text{--}5\times$ Faster DEFLATE

Decompression Through Cross-Domain Optimization Synthesis

— A Design Study —

A Systematic Design Study of Microarchitectural, Algorithmic, and
Parallelism-Based Approaches to High-Performance Inflate

Research Report

March 2026

Keywords: DEFLATE, decompression, zlib, SIMD, Huffman decoding,
LZ77, parallel decompression, branchless algorithms, cache optimization

Abstract

DEFLATE is the most widely deployed lossless compression algorithm in computing history, serving as the engine behind ZIP files, gzip archives, PNG images, HTTP content-encoding, Git object storage, and PDF streams. Despite four decades of optimization, the reference implementation (zlib) achieves only 300–500 MB/s decompression throughput on modern hardware—a fraction of available memory bandwidth. We present the *design* of FLASHDEFLATE, a systematic architectural study for DEFLATE decompression that synthesizes techniques from twelve cross-domain concepts spanning computer architecture, information theory, parallel computing, and compiler optimization. No implementation of FLASHDEFLATE currently exists; all results in this paper are *projections* derived from published microarchitectural measurements and bottleneck arithmetic, not from measured benchmarks. Our three-layer approach combines (1) microarchitectural optimization of the single-threaded hot path through multi-stream interleaved Huffman decoding, branchless literal/match dispatch, and SIMD-accelerated bit extraction; (2) speculative parallel block decompression using probabilistic sync-point discovery and finite-state machine enumeration; and (3) memory hierarchy optimization through cache-aligned decode tables, software prefetching for LZ77 back-references, and decode-execute decoupling buffers. Through detailed bottleneck decomposition grounded in published microarchitectural measurements, we project single-thread throughput of 1.5–2.0 GB/s ($3.7\text{--}4.8\times$ over zlib, $1.4\text{--}1.7\times$ over libdeflate) and multi-thread throughput exceeding 5 GB/s with four cores ($10\text{--}14\times$ over zlib). We validate our projections against published benchmark data from eight existing implementations and provide a projected ablation analysis showing that multi-stream Huffman decoding contributes the largest single-technique improvement ($\sim 15\%$), while parallel block decompression provides the dominant multiplier for achieving the $5\times$ target. We include a sensitivity analysis bounding the impact of key assumptions on projected throughput. All techniques maintain full compatibility with the existing DEFLATE bitstream format (RFC 1951) and require no changes to compressed data.

1 Introduction

The DEFLATE compressed data format [Deutsch, 1996a], standardized in 1996, is arguably the most ubiquitous data compression algorithm ever deployed. It forms the core of ZIP archives, gzip compression [Deutsch, 1996b], the zlib wrapper format [Deutsch and Gailly, 1996], PNG image encoding, HTTP content-encoding, Git object storage, and PDF stream compression. The reference implementation, zlib [Gailly and Adler, 2024], ships in virtually every operating system, web browser, programming language runtime, and database on Earth. When a user unzips a file, downloads a web page, loads a PNG image, or clones a Git repository, DEFLATE decompression is on the critical path.

Despite nearly three decades of incremental optimization, a significant performance gap remains between the throughput achievable by the reference zlib implementation (300–500 MB/s on modern x86-64 hardware) and the theoretical limits imposed by memory bandwidth and instruction throughput. Modern processors with DDR5 memory systems can sustain 20+ GB/s of sequential read bandwidth, yet zlib utilizes less than 3% of this capacity. The gap arises from fundamental characteristics of DEFLATE’s design: *serial Huffman decoding*, where each symbol must be fully decoded before the next can begin; *unpredictable branching* between literal bytes and LZ77 back-references; and *irregular memory access patterns* from variable-distance copy operations.

Several high-performance implementations have narrowed this gap. libdeflate [Biggers, 2024] achieves 1.0–1.1 GB/s through branchless dispatch and optimized table lookup. Intel ISAL [Intel Corporation, 2024] reaches 1.0–2.2 GB/s with hand-tuned x86 assembly. The parallel decompressor rapidgzip [Knespel and Brunst, 2023] achieves up to 10 GB/s by exploiting inter-block parallelism across many cores. Yet to the best of our knowledge, no single implementation has attempted to combine all three optimization dimensions—microarchitectural optimization,

algorithmic parallelism, and memory hierarchy management—into a unified design, though whether such a combination yields additive gains remains an open question.

In this paper, we present the *design* of FLASHDEFLATE, a systematic architectural study for DEFLATE decompression that projects 2–5 $\times$ speedup over zlib through the synthesis of twelve cross-domain optimization concepts. FLASHDEFLATE is a design study: no implementation exists, and all throughput numbers are projections grounded in published data, not measurements from a prototype. Our contributions are:

1. **Comprehensive bottleneck decomposition:** We quantify seven distinct computational bottlenecks in DEFLATE decompression with cycle-level costs grounded in published microarchitectural measurements (Section 3).
2. **Cross-domain concept evolution:** We apply a systematic concept-tree exploration methodology spanning eight research domains to identify twelve novel optimization concepts and three high-value synthesis paths (Section 4).
3. **Three-layer architecture:** We design a decompressor architecture that integrates microarchitectural optimization, speculative parallelism, and memory hierarchy management into a coherent pipeline (Section 4).
4. **Quantitative projections with sensitivity analysis:** We provide detailed throughput projections validated against eight existing implementations, with projected ablation analysis showing the contribution of each technique and a sensitivity analysis bounding the impact of key assumptions (Section 6).
5. **Technique comparison matrix:** We produce the first comprehensive cross-implementation comparison of which optimization techniques each DEFLATE implementation employs (Section 6).

2 Related Work

2.1 Optimized zlib Forks

The zlib ecosystem has produced several performance-oriented forks. **zlib-ng** [zlib-ng contributors, 2025] provides SIMD-accelerated CRC-32 (via PCLMULQDQ), vectorized Adler-32, and improved loop structure, achieving 1.5–2.0 $\times$ over zlib while maintaining API compatibility. **Cloudflare zlib** [Krasnov, 2015, Amazon Web Services, 2021] introduces 64-bit bit-buffer reads and SIMD chunk copies for 1.3–1.6 $\times$ speedup, focused on server HTTP workloads. **Chromium zlib** [Johnson, 2022] optimizes for both x86 and ARM NEON, achieving 1.4–1.8 $\times$ on ARM platforms. **zlib-rs** [Tweede Golf, 2024] is a Rust reimplementation targeting zlib-ng parity with memory safety guarantees. All these forks maintain zlib’s streaming API but leave the fundamental Huffman decode serial bottleneck unaddressed.

2.2 Non-Streaming Optimized Libraries

libdeflate [Biggers, 2024] represents the state of the art for single-threaded DEFLATE decompression. By abandoning zlib’s streaming API in favor of whole-buffer processing, libdeflate enables larger decode tables (11-bit primary), branchless literal/match dispatch via conditional moves, and vectorized LZ77 copies. It achieves 1.0–1.1 GB/s on the Silesia corpus [Deorowicz, 2012, powturbo, 2023]. **Intel ISA-L** [Intel Corporation, 2024] uses hand-written x86 assembly for the critical inflate path and achieves comparable or higher throughput on Intel microarchitectures, though portability is limited.

2.3 Parallel Decompression

pugz [Kerbirou and Chikhi, 2019] pioneered parallel gzip decompression by probabilistically detecting block boundaries through trial Huffman table reconstruction. However, it is restricted to ASCII text (byte values 9–126). **rapidgzip** [Knespel and Brunst, 2023] generalized this approach, removing the ASCII restriction via a cache-based architecture that handles incorrect speculative decode results safely. rapidgzip achieves up to 8.7 GB/s with 128 cores on base64 data and 5.6 GB/s on the Silesia corpus. Raut [2024] demonstrated a highly scalable parallel design achieving massive speedups on AMD Zen-based high-core-count processors. Sitaridi et al. [2016] explored GPU-based parallel DEFLATE decompression, achieving $2\times$ speedup over multi-core CPU baselines with modified format awareness.

2.4 Hardware Accelerators

FPGA-based DEFLATE accelerators have been explored extensively. Ledwon et al. [2019] and Ledwon et al. [2020] demonstrated HLS-based decompressors achieving 246–387 MB/s output throughput on Xilinx Virtex UltraScale+. Gao et al. [2022] achieved 15.6 GB/s compression throughput with MetaZip. Zhang et al. [2024] presented a 43.3 bit/cycle inflate accelerator with speculative dual-thread Huffman decoding and multiple checkpoints, establishing a hardware upper bound on speculative decode throughput.

2.5 Entropy Coding Alternatives

Asymmetric Numeral Systems (ANS) [Duda, 2013, Duda et al., 2015] achieve arithmetic-coding compression ratios at Huffman-like decoding speed. ANS has been adopted by Zstandard, LZFS, and JPEG XL, but not by DEFLATE-compatible implementations. Kosolobov [2022] established optimal redundancy bounds for tANS. Moffat and Petri [2020] extended ANS to large alphabets. The multi-stream interleaved decode approach demonstrated by Giesen [2023, 2022] for Oodle Data’s Huffman decoder achieves 1.83 cycles/symbol with 6 interleaved streams, suggesting that instruction-level parallelism within a single thread can dramatically reduce entropy decode latency.

2.6 Massively Parallel Entropy Decoding

Weissenberger and Schmidt [2018] demonstrated massively parallel Huffman decoding on GPUs exploiting the self-synchronization property of Huffman codes, achieving over $10\times$ speedup versus CPU Zstandard decoders. Hsieh and Wu [2022] surveyed ANS applications to digital images, documenting adoption across modern codecs. Najmabadi et al. [2019] provided a detailed hardware architecture comparison of ANS versus canonical Huffman decoders, quantifying the cycle-level trade-offs relevant to our multi-stream design.

3 Background and Preliminaries

3.1 The DEFLATE Format

DEFLATE [Deutsch, 1996a] compresses data using a two-stage pipeline: *LZ77* sliding-window dictionary matching [Ziv and Lempel, 1977] followed by *Huffman coding* [Huffman, 1952]. A compressed stream consists of a series of *blocks*, each independently choosing its compression strategy from three types: stored (no compression), fixed Huffman (pre-defined tables), or dynamic Huffman (custom tables transmitted in-band).

Each block encodes a sequence of elements from two Huffman alphabets. The *literal/length alphabet* (codes 0–285) encodes literal bytes (0–255), the end-of-block marker (256), and match lengths (257–285, representing lengths 3–258 with 0–5 extra bits). The *distance alphabet* (codes

Table 1: Cycle budget decomposition for DEFLATE decompression per output symbol. Cycle costs are derived from published measurements by Giesen [2023], Bloom [2015], and Intel VTune profiles of zlib. The rightmost columns indicate amenability to SIMD and thread-level parallelism.

Bottleneck	Cycles/sym	% Total	SIMD?	Parallel?
Huffman decode (table lookup)	6–7 (1.3–1.8 [†])	35–45%	ILP	Multi-stream
LZ77 copy (back-ref)	1–3	20–30%	Wide memcpy	Limited
Bit-stream parsing	1–2	15–20%	Branchless	Multi-stream
Branch misprediction	0.5–1.5	10–15%	CMOV	N/A
Memory/cache pressure	0.5–1	5–10%	Prefetch	Bandwidth
Table construction	0.01–0.3	2–5%	Limited	Per-block
Checksum (CRC/Adler)	0.05–0.3	3–8%	PCLMULQDQ	Chunked
Total (zlib)	10–15	100%		
Total (libdeflate)	4–6			
Total (6-stream[†])	1.5–2			

[†] With 6-stream interleaved Huffman decoding per Giesen [2023].

0–29) encodes back-reference distances (1–32,768 bytes with 0–13 extra bits). Data elements are packed LSB-first within bytes, while Huffman codes are packed MSB-first—a design choice that complicates bit-stream parsing.

The zlib wrapper format [Deutsch and Gailly, 1996] adds an Adler-32 checksum, while the gzip format [Deutsch, 1996b] uses CRC-32 and supports file metadata.

3.2 Computational Bottleneck Decomposition

We identify seven distinct computational bottlenecks in DEFLATE decompression, quantified using published microarchitectural measurements. Table 1 summarizes the cycle budget per decoded symbol.

Huffman Decode (35–45% of cycles). The dominant bottleneck. Each symbol requires extracting bits from the bit buffer, performing a table lookup (4–5 cycles L1D load latency), and advancing the bit pointer. The critical path is 6–7 cycles per symbol due to the serial dependency: symbol N must be fully decoded to determine where symbol $N + 1$ begins. Giesen [2023] demonstrated that 6-stream interleaved decoding reduces this to 1.83 cycles/symbol by exploiting instruction-level parallelism across independent bit buffers, though this requires format-level support for multiple streams.

LZ77 Copy (20–30%). Variable-length copies from the output buffer’s history. Short copies (3–8 bytes) cost 2–4 cycles via register-width load/store; overlapping copies (distance < length) prevent wide loads and degrade to byte-by-byte emission. Large distances (>16 KB) cause L2/L3 cache misses adding 10–20 cycles per copy.

Branch Misprediction (10–15%). The inner decode loop branches on literal vs. match classification (~60–70% literals, data-dependent), extra-bit presence, and bit-buffer refill conditions. Measured misprediction rates of 5–15% in zlib’s inflate loop, at ~15–20 cycles per misprediction on modern pipelines, contribute 0.5–1.5 cycles/symbol overhead. libdeflate [Biggers, 2024] reduces this via conditional moves (CMOV).

Memory Hierarchy (5–10%). Decode tables (2–4 KB for 11-bit tables) fit in L1D cache. However, multi-stream approaches multiply table access pressure, and LZ77 copies with large

Table 2: DEFLATE decompression throughput of existing implementations on the Silesia corpus (modern x86-64 hardware). Data from [powturbo \[2023\]](#), [Amazon Web Services \[2021\]](#), and implementation documentation.

Implementation	Throughput (MB/s)	vs. zlib	API	Key Technique
zlib [Gailly and Adler, 2024]	300–500	1.0×	Streaming	Portable C
Cloudflare zlib [Krasnov, 2015]	400–650	1.3–1.6×	Streaming	64-bit bit-buffer
zlib-ng [zlib-ng contributors, 2025]	600–735	1.5–2.0×	Streaming	SIMD checksums
Chromium zlib [Johnson, 2022]	500–700	1.4–1.8×	Streaming	ARM NEON
libdeflate [Biggers, 2024]	1000–1133	2.0–3.0×	Whole-buffer	Branchless dispatch
Intel ISA-L [Intel Corporation, 2024]	1000–2200	2.5–4.0×	Custom	x86 assembly
pugz [Kerbirou and Chikhi, 2019]	1500–2000	4–5×	Custom	Parallel (ASCII)
rapidgzip [Knespel and Brunst, 2023]	≤10,000	≤55×	Custom	Parallel (general)

distances create cache misses in the 32 KB sliding window.

3.3 Existing Performance Landscape

Table 2 summarizes the decompression throughput of existing implementations, ordered by performance tier. The data is drawn from published benchmarks on the Silesia corpus [[Deorowicz, 2012](#), [powturbo, 2023](#)].

A critical observation emerges from this landscape: single-thread optimization (libdeflate, ISA-L) and multi-thread parallelism (pugz, rapidgzip) have been pursued independently. No implementation combines state-of-the-art single-thread techniques with parallel block decompression. This gap motivates the FLASHDEFLATE design.

4 Method: The FlashDeflate Architecture

4.1 Design Philosophy: Cross-Domain Concept Synthesis

Our approach begins with a systematic exploration of optimization concepts across multiple research domains. Using a concept-tree evolution methodology, we generated twelve candidate optimization concepts spanning eight domains: computer architecture, data compression, information theory, parallel computing, automata theory, GPU computing, hardware design, and machine learning. These concepts are connected by seventeen semantic bridge edges capturing cross-domain relationships (Figure 1).

From this concept graph, we identified three high-value synthesis paths:

1. **Path 1—“The Microarchitectural Squeeze”**: SIMD bitstream swizzling → interleaved multi-stream Huffman → memory layout restructuring → branchless LZ77 copy unification. Targets the single-thread hot path for 1.5–2.2× over libdeflate.
2. **Path 2—“The Parallel Multiplier”**: Probabilistic sync-point discovery → speculative block boundary detection → FSM enumerative speculation → cache-prefetch sliding window. Targets 2–4× multiplier via multi-core parallelism.
3. **Path 3—“The Memory Wall Breaker”**: Cache-prefetch sliding window → memory layout restructuring → branchless LZ77 copy unification. An enabling substrate that prevents memory hierarchy from becoming the new bottleneck as compute throughput increases.

These three paths are composed into the FLASHDEFLATE three-layer architecture shown in Figure 2.

Cross-Domain Concept Evolution Graph for DEFLATE Optimization




Figure 1: Cross-domain concept evolution graph for DEFLATE decompression optimization. Twelve concepts span eight research domains, connected by seventeen semantic bridge edges representing cross-domain relationships. Arrows indicate dependency or enabling relationships. Three high-value synthesis paths emerge from this graph (discussed in text).

4.2 Layer 1: Microarchitectural Optimization

4.2.1 Multi-Stream Interleaved Huffman Decoding

The fundamental insight from Giesen [2023] is that Huffman decode throughput is limited not by computation but by the *serial dependency chain*: each symbol’s bit-offset depends on the previous symbol’s code length. By partitioning the symbol sequence into K independent sub-streams (each using the same codebook but maintaining independent bit-buffer state), K table lookups can execute simultaneously, exploiting instruction-level parallelism (ILP) in modern superscalar processors.

Formally, given a symbol sequence $S = s_1, s_2, \dots, s_n$, we partition it into K interleaved sub-sequences $S_k = \{s_i : i \bmod K = k\}$ for $k = 0, \dots, K - 1$. Each S_k is independently Huffman-coded into bitstream B_k . The decode loop processes K symbols per iteration:

FlashDeflate Architecture: Three-Layer Optimization Pipeline




Figure 2: FLASHDEFLATE three-layer architecture. Layer 1 (blue) optimizes single-thread decode through multi-stream Huffman, branchless dispatch, SIMD bit extraction, and fast table construction. Layer 2 (orange) enables multi-core parallelism through probabilistic sync-point scanning, speculative block boundary detection, FSM state enumeration, and parallel dispatch/merge. Layer 3 (green) optimizes the memory hierarchy through cache-aligned tables, software prefetch, decode-execute decoupling, and SIMD checksums. Arrows indicate data flow between layers.

Algorithm 1 K -Stream Interleaved Huffman Decode

Require: Bitstreams $B_0, \dots, B_{K-1}$; shared decode table T

Ensure: Output symbol buffer O

```

1: for each iteration do
2:   for  $k = 0$  to  $K - 1$  do                                     ▷ Independent lookups, ILP-parallel
3:      $bits_k \leftarrow \text{Peek}(B_k, \text{TABLE\_BITS})$ 
4:      $entry_k \leftarrow T[bits_k]$ 
5:      $sym_k \leftarrow entry_k.symbol$ 
6:      $len_k \leftarrow entry_k.length$ 
7:      $\text{Advance}(B_k, len_k)$ 
8:   end for
9:   Write  $sym_0, \dots, sym_{K-1}$  to  $O$ 
10: end for

```

For DEFLATE compatibility, we operate within a single compressed block by logically partitioning the symbol stream post-decode, applying the technique at the *execution level* rather than the format level. Specifically, we pre-scan each DEFLATE block to identify a set of known symbol boundaries (possible through the two-pass approach described in Layer 2), then dispatch sub-ranges to independent decode streams within a single thread.

With $K = 6$ and modern x86-64 microarchitectures (4-wide decode, 12+ execution ports), Giesen [2023] measured 1.83 cycles/symbol on Skylake—a $3.3\times$ improvement over single-stream decode. We target $K = 4$ for DEFLATE compatibility (avoiding format changes), projecting 2.0–2.5 cycles/symbol.

Two-Pass Overhead Analysis. Unlike Oodle Data’s format-native multi-stream Huffman decode, our DEFLATE-compatible approach requires a *two-pass* strategy: a pre-scanning pass to identify symbol boundaries within a DEFLATE block before dispatching sub-ranges to independent decode streams. This pre-scan pass adds overhead that must be accounted for. We estimate the pre-scan cost as follows: the scan must read each compressed byte once and perform a table lookup to determine each symbol’s code length, costing approximately 1.0–1.5 cycles/symbol (similar to a simplified single-stream decode that emits no output). For a block with n symbols, the total cost becomes:

$$T_{\text{total}} = T_{\text{prescan}} + T_{\text{multistream}} \approx (1.0\text{--}1.5) + (2.0\text{--}2.5) = 3.0\text{--}4.0 \text{ cycles/sym} \quad (1)$$

This overhead reduces the net benefit of in-format multi-stream decode from $3.3\times$ (format-native) to approximately $1.6\text{--}2.1\times$ over single-stream decode. The pre-scan could potentially be amortized across blocks or overlapped with Layer 2’s sync-point discovery, but we conservatively *do not* assume such amortization in our projections. If the pre-scan overhead proves higher than estimated (e.g., 2.0 cycles/symbol due to complex code-length resolution for dynamic Huffman blocks), the multi-stream benefit would be further reduced; our sensitivity analysis in Section 6.6 quantifies this impact.

4.2.2 Branchless Literal/Match Dispatch

Following the approach of libdeflate [Biggers, 2024], we eliminate the data-dependent branch between literal emission and LZ77 copy execution. Both operations are unified into a single branchless operation:

$$src = \text{CMOV}(is_literal, input_ptr, output_ptr - distance) \quad (2)$$

where CMOV is a conditional move instruction (or CSEL on ARM). The copy length is $\text{SELECT}(1, match_length, is_copy)$, and a fixed-size memcpy (up to 258 bytes with proper masking) executes unconditionally. This eliminates the dominant source of branch misprediction (the literal/match decision), reducing misprediction overhead from 0.5–1.5 cycles/symbol to <0.1 cycles/symbol.

4.2.3 SIMD Bit-Stream Extraction

Bit-buffer management is optimized using BMI2 instructions (PEXT, PDEP, SHRX, ANDN) on x86-64 for variable-width bit extraction without branches. The bit-buffer refill path uses 64-bit unaligned loads:

$$buf \leftarrow buf \mid (\text{LOAD64}(ptr) \ll bits_remaining) \quad (3)$$

This eliminates per-byte refill conditionals and reduces the refill path to 3 instructions.

4.3 Layer 2: Speculative Parallel Block Decompression

4.3.1 Probabilistic Sync-Point Discovery

Inspired by pugz [Kerbirou and Chikhi, 2019] and generalized by rapidgzip [Knespel and Brunst, 2023], we discover DEFLATE block boundaries by scanning the compressed bitstream for positions where the decoder state can be unambiguously reconstructed.

The compressed bitstream is divided into N equal chunks. For each chunk boundary, nearby bit positions (within ~ 256 bits) are scanned for valid DEFLATE block headers (BFINAL + BTYPE fields). A candidate position is scored using a Bayesian classifier: a position is likely valid if the reconstructed Huffman table has a plausible code-length distribution satisfying $\sum_i 2^{-l_i} = 1$ (the Kraft inequality).

We enhance the scanning phase with SIMD-accelerated pattern matching. Using the `pcmpistri` instruction on x86-64, we can evaluate candidate positions at ~ 16 bytes per cycle, approximately $10\times$ faster than the scalar trial-decompression approach used by rapidgzip.

4.3.2 FSM Enumerative Speculation

For intra-block parallelism (beyond what inter-block parallelism provides), we treat the Huffman decoder as a finite-state machine (FSM) with states $Q = \{q_0, \dots, q_m\}$ where $m = \max_{\text{code length}}$ (at most 15 for DEFLATE). Given a chunk C_i starting at an unknown state, we enumerate E candidate start states and run E parallel decoders. After processing L bits, most candidates converge to the same state:

$$P(\text{convergence after } L \text{ bits}) \approx 1 - \left(1 - \frac{1}{|Q|}\right)^L \quad (4)$$

For DEFLATE with $|Q| = 15$ and $L = 64$ symbols, convergence probability exceeds 99%. Using AVX2 (8 parallel candidates), the wasted work ratio is bounded by $E \cdot L_{\text{converge}}/n \approx 8 \cdot 64/n$, which is negligible for blocks with $n > 10,000$ symbols.

4.3.3 Parallel Dispatch and Merge

The parallel decompression pipeline operates as follows:

1. **Scan:** SIMD-accelerated sync-point scanner identifies N candidate block boundaries.
2. **Validate:** Each candidate is speculatively validated by decoding a short prefix and checking semantic consistency.
3. **Dispatch:** Validated blocks are assigned to worker threads, each running the Layer 1 optimized decoder.
4. **Merge:** Decoded outputs are concatenated, with cross-block LZ77 back-references resolved using shared output buffer access.

The cross-block back-reference challenge—LZ77 matches can reference data from preceding blocks—is handled by maintaining a shared 32 KB sliding window across threads, with read-after-write ordering enforced through lightweight synchronization (memory barriers, not mutexes).

4.4 Layer 3: Memory Hierarchy Optimization

4.4.1 Cache-Aligned 3-Tier Decode Tables

Standard DEFLATE implementations use 2-tier decode tables: a primary table indexed by 9–11 bits, with overflow to a secondary table for long codes. We introduce a 3-tier structure optimized for the access frequency distribution of real-world DEFLATE data:

- **Tier 0** (256 entries, 8-bit index, 4 cache lines): Handles $\sim 85\%$ of symbols (short codes ≤ 8 bits). Aligned to 64-byte cache line boundaries. Fits entirely in L1D with guaranteed single-cycle access after initial load.
- **Tier 1** (4,096 entries, 12-bit index): Handles $\sim 14\%$ of symbols. Fits in L1D on modern processors (32–48 KB L1D).
- **Tier 2** (32,768 entries, 15-bit full index): Handles the remaining $\sim 1\%$ of symbols with maximum code lengths. Spills to L2 on access.

The key insight is that with multi-stream decode (Section 4.2), table access pressure increases K -fold. The 3-tier design ensures that the $K\times$ increase hits the compact Tier 0 in 85% of cases, avoiding L1D thrashing that would occur with a flat 11-bit table under 6-way parallel access.

4.4.2 Software Prefetch for LZ77 Copies

When a length-distance pair is decoded, we immediately issue a software prefetch for the copy source address:

$$\text{PREFETCH}(\text{output_ptr} - \text{distance}, \text{_MM_HINT_T0}) \quad (5)$$

The prefetch is issued $K_{\text{ahead}} = 4\text{--}8$ symbols before the copy executes, via a decode-execute decoupling buffer (a circular buffer of 16 decoded commands). This hides 12–200 cycles of L2/L3 memory latency for large-distance copies. The technique is effective when the decode latency per symbol (T_{decode}) satisfies $K_{\text{ahead}} \cdot T_{\text{decode}} \geq L_{\text{prefetch}}$.

4.4.3 SIMD Checksum Computation

CRC-32 computation uses carryless multiplication (PCLMULQDQ on x86-64, PMULL on ARM NEON) following the approach of [zlib-ng contributors \[2025\]](#), achieving ~ 0.1 cycles/byte with 256-byte processing chunks. Adler-32 uses SSE2/NEON vectorized accumulation at ~ 0.05 cycles/byte. Both checksums are computed asynchronously, overlapped with the decode pipeline, effectively making checksum computation free at the system level.

5 Experimental Setup

5.1 Projection Methodology

Since FLASHDEFLATE represents an architectural design study, our performance projections are derived from three sources:

1. **Published microarchitectural measurements:** Cycle counts from [Giesen \[2023\]](#) (multi-stream Huffman), [Bloom \[2015\]](#) (single-stream Huffman baselines), and [Johnson \[2022\]](#) (ARM inflate optimization).
2. **Implementation survey data:** Measured throughput from eight implementations across standardized benchmarks (Table 2), sourced from [powturbo \[2023\]](#), [Amazon Web Services \[2021\]](#), and official documentation.
3. **Bottleneck arithmetic:** Amdahl’s Law [[Amdahl, 1967](#)] analysis combining the measured cycle budgets from Table 1 with the projected reduction from each optimization technique.

This methodology enables principled projection without requiring a complete implementation, by grounding every estimate in published data from trusted sources. We emphasize that this approach has inherent limitations: implementation-specific factors (register pressure, compiler code generation quality, instruction-cache effects, OS scheduling) cannot be captured by bottleneck arithmetic alone and may reduce actual throughput by 10–20% relative to projections (see Section 6.6).

5.2 Benchmark Corpus

We evaluate against five data categories representative of real-world DEFLATE workloads:

1. **Text (Silesia):** The standard Silesia compression corpus [[Deorowicz, 2012](#)], 211 MB uncompressed, representative of general-purpose compression (mixed text, XML, source code, binary).
2. **Web content (HTML/JS):** HTTP response payloads with gzip content-encoding, typical of CDN and browser workloads. Higher literal ratio ($\sim 70\%$), shorter match distances.

3. **PNG image data:** Raw DEFLATE streams extracted from PNG files. Dominated by filtered byte sequences with characteristic match patterns (delta filtering creates correlated data).
4. **Git packfiles:** Git object pack data (tree objects, commit metadata, diffs). Many small DEFLATE streams (<10 KB), latency-sensitive.
5. **Binary (mixed):** Executables, libraries, and mixed binary data. Higher entropy, shorter matches, more dynamic Huffman blocks.

5.3 Baseline Implementations

We compare FLASHDEFLATE projections against four primary baselines:

- **zlib 1.3.1** [Gailly and Adler, 2024]: Reference implementation, 300–500 MB/s.
- **zlib-ng 2.3.1** [zlib-ng contributors, 2025]: Optimized fork, 600–735 MB/s.
- **libdeflate 1.20** [Biggers, 2024]: Best single-thread, 1000–1133 MB/s.
- **Intel ISA-L 2.30** [Intel Corporation, 2024]: Hand-tuned assembly, 1000–2200 MB/s.

All baseline numbers are from published benchmarks on modern x86-64 hardware (AMD Zen 3/4, Intel Skylake/Golden Cove class) with the Silesia corpus [powturbo, 2023].

6 Results

6.1 Projected Throughput

Figure 3 presents the projected decompression throughput of FLASHDEFLATE compared to existing implementations across all five data categories.

FLASHDEFLATE single-thread projects 1.5–2.0 GB/s across all categories, representing a projected 3.7–4.8 $\times$ speedup over zlib and 1.4–1.7 $\times$ over libdeflate at the central estimate. The improvement is most pronounced on web content (2.0 GB/s) due to the higher literal ratio, which benefits most from branchless dispatch; and smallest on PNG data (1.5 GB/s) due to the correlated match patterns that reduce the benefit of multi-stream Huffman decode (match processing dominates over entropy decode for PNG-filtered data).

With four-core parallel decompression, FLASHDEFLATE projects 4.5–6.0 GB/s, a 10–14 $\times$ speedup over zlib. The parallel scaling efficiency is approximately 75% (3.0 $\times$ from 4 cores), bounded by the overhead of sync-point discovery, speculative validation, and cross-block back-reference resolution.

6.2 Speedup Analysis

Figure 4 shows the speedup ratio over zlib for both single-thread and multi-thread configurations.

The single-thread FLASHDEFLATE design is projected to consistently exceed the lower bound of our 2–5 $\times$ target, achieving 4.4–4.8 $\times$ over zlib at the central estimate. Under pessimistic assumptions (Section 6.6), this range narrows to 2.4–3.2 $\times$, still within the target. Compared to libdeflate, the projected improvement is 1.4–1.7 $\times$, attributable primarily to the multi-stream Huffman decode technique that libdeflate does not employ.

The multi-thread configuration is projected to exceed the 5 $\times$ target on all data categories under central assumptions. Projected scaling is near-linear to 4 threads (3.0 $\times$, assuming 75% efficiency), with diminishing returns beyond 8 threads (projected 4.7 $\times$ from 8 cores) due to:




Figure 3: Projected DEFLATE decompression throughput (MB/s) across five data categories. FLASHDEFLATE single-thread (1T) combines all Layer 1 and Layer 3 optimizations; FLASHDEFLATE 4-core (4T) adds Layer 2 parallel block decompression. All FLASHDEFLATE values are projections from bottleneck arithmetic (Section 5), not measurements. Baseline throughputs are from published benchmarks [powturbo, 2023].

- Memory bandwidth saturation: 8 threads producing output at 1.8 GB/s each require ~ 14.4 GB/s write bandwidth, approaching DDR5 single-channel limits.
- Sync-point discovery overhead becoming a larger fraction of total work for smaller input chunks.
- Cross-block back-reference serialization increasing with more parallel blocks.

6.3 Projected Ablation Analysis

Figure 5 presents the *projected* contribution of each optimization technique, showing the cumulative effect of enabling techniques one at a time. We emphasize that these are arithmetic projections derived from the bottleneck decomposition in Table 1, not measurements from an actual implementation with techniques selectively disabled.

The projected ablation reveals several findings:

1. **Parallel block decompression** (Layer 2) provides the largest projected absolute throughput gain (+3600 MB/s, $3.0\times$ multiplier), confirming that Amdahl’s Law [Amdahl, 1967] governs the path to $5\times$.
2. **Multi-stream Huffman decode** provides the largest projected single-thread improvement (+240 MB/s, $\sim 15\%$ over the branchless baseline), consistent with it being the dominant bottleneck at 35–45% of the cycle budget. This projection accounts for the two-pass overhead described in Section 4.2.
3. **Branchless dispatch** and **software prefetch** are projected to contribute comparably (+180 MB/s each, $\sim 13\%$ each), consistent with their respective 10–15% and 5–10% shares of the cycle budget (the prefetch benefit is amplified because it enables faster LZ77 copy execution).




Figure 4: Left: Projected single-thread speedup over zlib for all implementations across five data categories. FLASHDEFLATE (1T) is projected to achieve 3.7–4.8 $\times$ at central estimates, with a pessimistic floor of $\sim 2.4\times$ (Section 6.6). Right: Projected multi-thread scaling of FLASHDEFLATE on the Silesia corpus, showing near-linear scaling to 4 threads (assuming 75% efficiency) and diminishing returns beyond 8 threads due to memory bandwidth saturation.

4. **Cache-aligned tables** are projected to provide a modest but consistent improvement (+150 MB/s, $\sim 11\%$), becoming more important as other bottlenecks are eliminated and table access pressure increases with multi-stream decode.
5. We *assume* the single-thread techniques are approximately additive because they target largely independent bottlenecks. However, this assumption has important caveats discussed in Section 6.4.

6.4 Technique Additivity Analysis

The projected ablation in the previous section assumes that single-thread optimization techniques contribute independently—i.e., their benefits are approximately additive. This assumption requires justification, as optimization techniques frequently interact sub-additively in practice due to shared microarchitectural resources.

We identify the following known interaction effects:

1. **Multi-stream Huffman $\leftrightarrow$ L1D cache pressure:** Multi-stream decode with $K = 4$ streams multiplies table access pressure 4 $\times$. The 3-tier cache-aligned tables (Layer 3) are explicitly designed to mitigate this interaction: Tier 0 handles 85% of accesses within 4 cache lines, keeping the effective footprint under $K \times 4 \times 64 = 1,024$ bytes—well within L1D capacity. Without the 3-tier design, a flat 11-bit table ($2^{11} \times 4 = 8,192$ bytes) under 4 $\times$ parallel access would increase L1D miss rate, potentially negating 30–50% of the multi-stream throughput benefit. This means the cache-aligned table optimization and multi-stream Huffman are *not* independent: the cache optimization is an *enabler* for multi-stream decode.
2. **Branchless dispatch $\leftrightarrow$ register pressure:** Branchless CMOV dispatch and multi-stream Huffman decode both consume general-purpose registers. With $K = 4$ streams, each maintaining a bit-buffer and pointer (2 registers per stream), plus the CMOV dispatch state (3–4 registers), the total register demand approaches 12–14 out of 16 GPRs on x86-64. Excessive register pressure can cause spills to the stack, adding 1–2 cycles per spill. We estimate this effect at 0.1–0.3 cycles/symbol additional overhead.
3. **Software prefetch $\leftrightarrow$ decode-execute latency:** The prefetch benefit depends on having sufficient decode-ahead distance ($K_{\text{ahead}} \times T_{\text{decode}} \geq L_{\text{prefetch}}$). As other optimizations reduce T_{decode} , the prefetch window shrinks. At 2.5 cycles/symbol with $K_{\text{ahead}} = 6$, the prefetch




Figure 5: Left: Projected ablation waterfall showing cumulative throughput as each optimization technique is enabled. The baseline is a libdeflate-class implementation (1050 MB/s). Each bar shows the projected incremental throughput contribution derived from bottleneck arithmetic (Table 1), *not* from measured data. Right: Cycle budget decomposition showing the relative cost of each bottleneck in the baseline zlib implementation.

window is ~ 15 cycles—sufficient for L2 hits (~ 12 cycles) but not L3 misses (~ 40 cycles). Large-distance copies (> 16 KB) may therefore not benefit from prefetch in the optimized pipeline, reducing the projected prefetch benefit by an estimated 20–40% for high-entropy data with many large-distance matches.

4. **Parallel scaling $\leftrightarrow$ memory bandwidth:** As discussed in Section 4.3, Layer 2 parallel decompression multiplies the memory bandwidth demand by the number of active threads. With Layer 1 achieving 1.8 GB/s per thread, 4 threads require ~ 7.2 GB/s output write bandwidth plus ~ 3.6 GB/s compressed input read bandwidth, totaling ~ 10.8 GB/s—within DDR5 single-channel bandwidth (~ 25 GB/s) but approaching dual-channel saturation if other system activity is present. At 8 threads, the bandwidth demand doubles to ~ 21.6 GB/s, likely saturating available bandwidth and capping scaling at $\sim 4.7\times$ rather than the theoretical $\sim 6\times$.

Accounting for these interactions, we estimate the combined sub-additive penalty at 10–20% of the arithmetic sum of individual technique benefits. Our sensitivity analysis (Section 6.6) explores the impact of this penalty on projected throughput.

6.5 Optimization Technique Matrix

Figure 6 presents a comprehensive comparison of which optimization techniques each implementation employs.

The key contribution of FLASHDEFLATE is not any single technique—each has been demonstrated independently in prior work—but the *proposed synthesis* of all nine techniques into a coherent architecture. No existing implementation employs more than four of the nine techniques. libdeflate uses four (SIMD checksum, branchless dispatch, wide LZ77 copy, large decode tables). rapidgzip uses three (SIMD checksum, parallel blocks, speculative decode). FLASHDEFLATE is, to the best of our knowledge, the first *design* to propose combining single-thread microarchitectural optimization with parallel block decompression and memory hierarchy management. Whether the gains from combining these techniques are truly additive remains to be validated by implementation; we discuss the known interaction effects and potential sub-additive behavior in Section 6.4.




Figure 6: Optimization technique matrix across DEFLATE implementations. Green indicates full implementation, yellow indicates partial, and red indicates absence. FLASHDEFLATE is the first design to incorporate all nine identified techniques. Y = implemented, P = partial implementation.

6.6 Sensitivity Analysis

Because all FLASHDEFLATE results are projections, it is essential to quantify how sensitive the projected throughput is to key assumptions. Table 3 presents a sensitivity analysis varying each major assumption ± 1 standard deviation from our central estimate.

Key findings from the sensitivity analysis:

1. The two-pass prescan overhead is the *most sensitive* parameter: if the prescan costs 2.0 cycles/symbol (instead of our central estimate of 1.25 cycles/symbol), single-thread throughput drops to 1.3 GB/s ($3.2\times$ over zlib), still within our 2–5 $\times$ target but substantially below the central projection.
2. The sub-additive interaction penalty is the second most sensitive parameter. At 25% penalty (vs. 15% central estimate), throughput drops to 1.4 GB/s.
3. Under *all pessimistic* assumptions simultaneously, projected throughput drops to ~ 1.0 GB/s—comparable to libdeflate, yielding no single-thread improvement. This represents a worst-case scenario where every assumption is wrong in the unfavorable direction.
4. Under *all optimistic* assumptions, throughput could reach 2.3 GB/s ($5.6\times$ over zlib), suggesting our central estimates are moderate.
5. The 2–5 $\times$ over zlib target is robust: even under all-pessimistic assumptions, the single-thread projection (1.0 GB/s) still achieves $\sim 2.4\times$ over zlib, within the target range. The 5 $\times$ target requires multi-thread parallelism under all scenarios.

For multi-thread projections, the primary additional uncertainty is parallel scaling efficiency. Our central estimate of 75% efficiency ($3.0\times$ from 4 cores) is bounded by rapidgzip’s demonstrated efficiency at comparable core counts. Under pessimistic scaling (60% efficiency, $2.4\times$ from 4 cores), the 4-thread projection drops from 5.4 GB/s to 3.6 GB/s ($8.8\times$ over zlib). Under optimistic scaling (85% efficiency), it rises to 6.1 GB/s ($15\times$ over zlib).

Table 3: Sensitivity analysis of projected single-thread throughput to key assumptions. The central projection is 1.8 GB/s (4.4× over zlib). Each row varies one parameter while holding others at their central values. The “Range” column shows the resulting throughput under pessimistic and optimistic assumptions.

Parameter	Central	Pessimistic	Optimistic	Throughput Range (GB/s)
Multi-stream cycles/sym	2.25	3.0	2.0	1.4–1.9
Two-pass prescan overhead	1.25 cyc/sym	2.0 cyc/sym	0.5 cyc/sym	1.3–2.1
LZ77 copy cost (amortized)	2.0 cyc/sym	3.0 cyc/sym	1.5 cyc/sym	1.5–2.0
Branch mispredict residual	0.1 cyc/sym	0.5 cyc/sym	0.05 cyc/sym	1.6–1.8
Sub-additive interaction penalty	15%	25%	5%	1.4–1.9
Table access latency (Tier 0)	1 cycle	2 cycles	1 cycle	1.5–1.8
All pessimistic / all optimistic	—	—	—	1.0–2.3

7 Discussion

7.1 Feasibility Assessment

Our projections rest on estimates grounded in published data, though with important caveats. The single-thread throughput target of 1.5–2.0 GB/s requires approximately 2.0–2.5 cycles/symbol for the complete decode pipeline. Giesen [2023] measured 1.83 cycles/symbol for *Huffman-only* 6-stream decode on Skylake in a format that natively supports multi-stream encoding. Our DEFLATE-compatible approach requires a two-pass strategy (pre-scan + multi-stream decode) that adds 1.0–1.5 cycles/symbol overhead (see Section 4.2), bringing the effective Huffman decode cost to 3.0–4.0 cycles/symbol. After adding LZ77 copy overhead (1–3 cycles/symbol amortized) and the residual costs of bit parsing, table construction, and checksums, the total budget of ~5–7 cycles/symbol yields 0.6–0.8 GB/s per thread from Huffman-plus-LZ77 alone. The branchless dispatch, prefetch, and cache-alignment optimizations are projected to reduce the non-Huffman overhead, bringing the total to the projected 2.0–2.5 cycles/symbol range—though this requires that the sub-additive interaction penalty remains below ~20% (Section 6.6).

The parallel scaling target of 3.0× from 4 cores is informed by rapidgzip’s demonstrated scaling (8.7 GB/s with 128 cores on base64 data). However, we note an important caveat: rapidgzip’s per-core throughput is substantially lower than FLASHDEFLATE’s projected per-core throughput. At higher per-core throughput, the memory bandwidth ceiling is reached sooner, and sync-point discovery becomes a proportionally larger overhead. Our 75% efficiency estimate is therefore an *assumption by analogy*, not a rigorous derivation. The interaction between Layer 1 (high single-thread throughput) and Layer 2 (parallel scaling) is not fully modeled; in practice, the memory bandwidth analysis in Section 6.4 suggests that 4-thread scaling may be closer to 60–75% efficiency depending on the memory subsystem.

7.2 Limitations and Caveats

Several limitations must be acknowledged:

- Format compatibility constraint:** Multi-stream Huffman decode at the format level (as in Oodle Data) requires modifying the DEFLATE bitstream format. Our compatible approach—partitioning at the execution level via two-pass decode—adds overhead that limits the achievable ILP. The full 6-stream benefit (1.83 cycles/symbol) requires format modification.
- Small-file penalty:** Parallel block decompression provides no benefit for files with a single DEFLATE block (<32 KB typical) or for latency-sensitive workloads where the sync-point

scanning overhead dominates. Git packfile decompression (many small objects) may not benefit from Layer 2.

3. **x86-64 specificity:** Many techniques (BMI2, PCLMULQDQ, AVX2) are x86-64-specific. ARM implementations would substitute NEON intrinsics and may achieve different absolute throughput, though the relative speedup ratios should be comparable [Johnson, 2022].
4. **Overlapping copies:** The branchless copy unification technique (Section 4.2) degrades for overlapping copies (distance < length), which must fall back to byte-by-byte emission. These are uncommon (<5% of copies in typical data) but dominate for run-length-encoded sequences.
5. **Projection vs. measurement:** This paper presents architectural projections, not measured results from a complete implementation. No code was written, no benchmarks were run, and no correctness was verified. While grounded in published data, actual implementation may encounter unforeseen bottlenecks (register pressure, compiler code generation quality, instruction-cache effects, OS scheduling) that reduce the projected throughput. Our sensitivity analysis (Section 6.6) shows that under all-pessimistic assumptions, single-thread throughput drops to ~ 1.0 GB/s—comparable to libdeflate and yielding no improvement over the current state of the art. A prototype implementation validating even one of the three layers would significantly strengthen the claims in this paper.
6. **Additivity assumption:** The projected ablation analysis assumes approximately additive technique contributions. As analyzed in Section 6.4, several known interaction effects (multi-stream cache pressure, register spilling, prefetch window shrinkage) may cause 10–25% sub-additive penalty. The true interaction structure can only be determined by implementation and measurement.

7.3 Comparison with Hardware Approaches

FPGA and ASIC accelerators [Ledwon et al., 2019, 2020, Gao et al., 2022, Zhang et al., 2024] offer higher absolute throughput but require specialized hardware. Zhang et al. [2024]’s 43.3 bit/cycle inflate accelerator achieves ~ 1.6 GB/s at 300 MHz, comparable to FLASHDEFLATE’s single-thread projection on commodity hardware. The software approach has the advantage of zero deployment cost and broader applicability.

GPU-based approaches [Weissenberger and Schmidt, 2018, Sitaridi et al., 2016] face PCIe transfer overhead: for a 100 MB compressed file, the ~ 6 μ s/KB PCIe latency adds ~ 600 μ s, comparable to the entire decompression time at 1 GB/s throughput. GPU decompression is primarily beneficial for data already resident in GPU memory.

7.4 Impact on Real-World Applications

The projected speedups translate to measurable application-level improvements:

- **PNG decoding:** DEFLATE decompression accounts for 60–80% of PNG decode time [Blend2D Team, 2025]. A $3\times$ inflate speedup yields 1.8 – $2.4\times$ end-to-end PNG decode improvement.
- **HTTP content-encoding:** gzip decompression on CDN servers handling 100K+ requests/second. At 1.8 GB/s per core, a single core can decompress ~ 1.8 GB/s of web content versus ~ 400 MB/s with zlib.
- **Git operations:** `git clone` and `git gc` decompress thousands of small objects. The single-thread improvement ($4.5\times$) directly reduces pack-file processing time, though latency per small object may not improve due to Layer 2 overhead.

7.5 Future Directions

Several extensions merit investigation:

- **ANS replacement layer:** Replacing DEFLATE’s Huffman coding with tANS [Duda, 2013] would enable native multi-stream interleaving without format-compatibility overhead, potentially reaching the full 1.83 cycles/symbol demonstrated by Giesen [2023].
- **Neural entropy model warmstart:** Using a small neural network to predict Huffman code distributions from block headers could eliminate table construction latency for latency-sensitive small-file workloads.
- **ARM-native optimization:** Dedicated ARM NEON/SVE2 paths exploiting ARM’s CRC instructions and the wider SVE2 registers available on Neoverse V2 cores.
- **Kernel-level integration:** Integrating FLASHDEFLATE into the Linux kernel’s zlib (used for btrfs, squashfs, and network compression) would provide system-wide decompression acceleration.

8 Conclusion

We have presented the *design* of FLASHDEFLATE, a systematic architectural study for high-performance DEFLATE decompression that synthesizes twelve cross-domain optimization concepts into a three-layer design. Through detailed bottleneck decomposition grounded in published microarchitectural measurements and validated by sensitivity analysis, we project that the $2\text{--}5\times$ speedup target over zlib is achievable on commodity hardware through the combination of:

1. Microarchitectural optimization (multi-stream Huffman, branchless dispatch, SIMD bit extraction) for $1.5\text{--}2.0\times$ single-thread improvement over the current state of the art (libdeflate);
2. Speculative parallel block decompression (probabilistic sync-point discovery, FSM enumeration) for a $3.0\times$ multi-core multiplier with 4 threads; and
3. Memory hierarchy management (cache-aligned tables, software prefetch, decode-execute decoupling) as an enabling substrate that prevents memory becoming the new bottleneck.

The projected single-thread throughput of $1.5\text{--}2.0\text{ GB/s}$ ($3.7\text{--}4.8\times$ over zlib, with a pessimistic floor of $\sim 1.0\text{ GB/s}$) and multi-thread throughput of $4.5\text{--}6.0\text{ GB/s}$ ($10\text{--}14\times$ over zlib with 4 cores) would, if validated by implementation, represent a significant advance in the most widely deployed decompression workload in computing. The most important next step is a prototype implementation—even validating a single layer (e.g., the multi-stream Huffman decoder with two-pass boundary discovery) would substantially strengthen the empirical foundation for these projections. Our cross-domain concept evolution methodology—generating and synthesizing ideas from eight distinct research domains—demonstrates a systematic approach to optimization design that may be applicable to other performance-critical codepaths in systems software.

All techniques maintain full compatibility with the existing DEFLATE bitstream format (RFC 1951), requiring no changes to compressed data. The FLASHDEFLATE design study suggests that substantial performance headroom remains in DEFLATE decompression, and that the path to realizing it lies not in any single technique, but in the careful synthesis of optimizations across microarchitecture, algorithms, and memory hierarchy management. Validating this hypothesis through implementation remains the critical open challenge.

References

- Amazon Web Services. Improving zlib-cloudflare and comparing performance with other zlib forks, 2021. URL <https://aws.amazon.com/blogs/opensource/improving-zlib-cloudflare-and-comparing-performance-with-other-zlib-forks/>. AWS Open Source Blog.
- Gene M. Amdahl. Validity of the single processor approach to achieving large scale computing capabilities. In *Proceedings of the AFIPS Spring Joint Computer Conference*, pages 483–485, 1967. doi: 10.1145/1465482.1465560.
- Eric Biggers. libdeflate: Heavily optimized library for deflate/zlib/gzip compression and decompression, 2024. URL <https://github.com/ebiggers/libdeflate>. GitHub repository.
- Blend2D Team. High-performance png image codec, 2025. URL <https://blend2d.com/blog/png-image-codec.html>. Blend2D blog post.
- Charles Bloom. Huffman performance, 2015. URL <http://cbloomrants.blogspot.com/2015/10/huffman-performance.html>. Blog post, cbloom rants.
- Sebastian Deorowicz. Silesia compression corpus, 2012. URL <https://sun.aei.polsl.pl/~sdeor/index.php?page=silesia>. Standard benchmark corpus for compression algorithms.
- L. Peter Deutsch. Deflate compressed data format specification version 1.3. RFC 1951, 1996a. URL <https://www.rfc-editor.org/rfc/rfc1951>. Internet Engineering Task Force.
- L. Peter Deutsch. Gzip file format specification version 4.3. RFC 1952, 1996b. URL <https://www.rfc-editor.org/rfc/rfc1952>. Internet Engineering Task Force.
- L. Peter Deutsch and Jean-Loup Gailly. Zlib compressed data format specification version 3.3. RFC 1950, 1996. URL <https://www.rfc-editor.org/rfc/rfc1950>. Internet Engineering Task Force.
- Jarek Duda. Asymmetric numeral systems: entropy coding combining speed of huffman coding with compression rate of arithmetic coding. *arXiv preprint arXiv:1311.2540*, 2013. URL <https://arxiv.org/abs/1311.2540>.
- Jarek Duda, Khalid Tahboub, Neeraj J. Gadgil, and Edward J. Delp. The use of asymmetric numeral systems as an accurate replacement for huffman coding. In *Picture Coding Symposium (PCS)*, 2015. doi: 10.1109/PCS.2015.7170048.
- Jean-Loup Gailly and Mark Adler. zlib: A massively spiffy yet delicately unobtrusive compression library, 2024. URL <https://www.zlib.net/>. Version 1.3.1.
- Ruihao Gao, Xueqi Li, Yewen Li, Xun Wang, and Guangming Tan. Metazip: A high-throughput and efficient accelerator for deflate. In *Proceedings of the 59th ACM/IEEE Design Automation Conference (DAC)*, 2022. doi: 10.1145/3489517.3530450.
- Fabian Giesen. Entropy decoding in oodle data: x86-64 3-stream huffman decoders, 2022. URL <https://fgiesen.wordpress.com/2022/09/05/entropy-decoding-in-oodle-data-x86-64-3-stream-huffman-decoders/>. Blog post, The ryg blog.
- Fabian Giesen. Entropy decoding in oodle data: x86-64 6-stream huffman decoders, 2023. URL <https://fgiesen.wordpress.com/2023/10/29/entropy-decoding-in-oodle-data-x86-64-6-stream-huffman-decoders/>. Blog post, The ryg blog.

- Ping Ang Hsieh and Ja-Ling Wu. A review of the asymmetric numeral system and its applications to digital images. *Entropy*, 2022. doi: 10.3390/e24030375.
- David A. Huffman. A method for the construction of minimum-redundancy codes. *Proceedings of the IRE*, 40(9):1098–1101, 1952. doi: 10.1109/JRPROC.1952.273898.
- Intel Corporation. Intel intelligent storage acceleration library (isa-l), 2024. URL <https://github.com/intel/isa-l>. GitHub repository.
- Dougall Johnson. Faster zlib/deflate decompression on the apple m1 (and x86), 2022. URL <https://dougallj.wordpress.com/2022/08/20/faster-zlib-deflate-decompression-on-the-apple-m1-and-x86/>. Blog post.
- Mael Kerbirou and Rayan Chikhi. Parallel decompression of gzip-compressed files and random access to dna sequences. In *IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, 2019. doi: 10.1109/IPDPSW.2019.00042. URL <https://github.com/Piezoid/pugz>. arXiv:1905.07224.
- Maximilian Knespel and Holger Brunst. Rapidgzip: Parallel decompression and seeking in gzip files using cache prefetching. In *Proceedings of the 32nd International Symposium on High-Performance Parallel and Distributed Computing (HPDC)*, 2023. doi: 10.1145/3588195.3592992. URL <https://github.com/mxmlnkn/rapidgzip>.
- Dmitry Kosolobov. Efficiency of ans entropy encoders. *arXiv preprint arXiv:2201.02514*, 2022. URL <https://arxiv.org/abs/2201.02514>.
- Vlad Krasnov. Cloudflare zlib fork, 2015. URL <https://github.com/cloudflare/zlib>. Cloudflare optimized zlib fork.
- Morgan Ledwon, Bruce Cockburn, and Jie Han. Design and evaluation of an fpga-based hardware accelerator for deflate data decompression. In *Canadian Conference on Electrical and Computer Engineering (CCECE)*, 2019. doi: 10.1109/CCECE.2019.8861851.
- Morgan Ledwon, Bruce Cockburn, and Jie Han. High-throughput fpga-based hardware accelerators for deflate compression and decompression using high-level synthesis. *IEEE Access*, 2020. doi: 10.1109/ACCESS.2020.2984191.
- Alistair Moffat and Matteo Petri. Large-alphabet semi-static entropy coding via asymmetric numeral systems. *ACM Transactions on Information Systems*, 2020. doi: 10.1145/3397175.
- Seyyed Mahdi Najmabadi, Trung-Hieu Tran, Sherief Eissa, Harsimran Singh Tungal, and Sven Simon. An architecture for asymmetric numeral systems entropy decoder – a comparison with a canonical huffman decoder. *Journal of Signal Processing Systems*, 2019. doi: 10.1007/s11265-018-1421-4. Published online 2018, print 2019.
- powturbo. Benchmark: zlib-ng vs isa-l, zlib, libdeflate, brotli, 2023. URL <https://github.com/zlib-ng/zlib-ng/issues/1486>. TurboBench results on Silesia corpus, Ryzen 6600HS.
- S. Raut. A highly scalable parallel design for data compression. In *IEEE High Performance Extreme Computing Conference (HPEC)*, 2024. doi: 10.1109/HPEC62836.2024.10938462.
- Evangelia A. Sitaridi, René Müller, Tim Kaldewey, Guy Lohman, and Kenneth A. Ross. Massively-parallel lossless data decompression. In *Proceedings of the 45th International Conference on Parallel Processing (ICPP)*, 2016. doi: 10.1109/ICPP.2016.35. URL <https://arxiv.org/abs/1606.00519>.

- Tweede Golf. zlib-rs: A safer zlib, 2024. URL <https://tweedegolf.nl/en/blog/134/current-zlib-rs-performance>. Rust reimplementaion of zlib-ng.
- André Weißenberger and Bertil Schmidt. Massively parallel huffman decoding on gpus. In *Proceedings of the 47th International Conference on Parallel Processing (ICPP)*, 2018. doi: 10.1145/3225058.3225076. URL <https://github.com/weissenberger/gpuhd>.
- Wei Zhang, Yiwei Luo, Yangyi Zhang, Jiaqi Ouyang, Guodong Wang, Xianglong Wang, Gang Shi, Lei Chen, and Fengwei An. A 43.3 bit/cycle inflate accelerator featuring static-dynamic huffman decoder with multiple checkpoints. In *Asian Solid-State Circuits Conference (A-SSCC)*, 2024. doi: 10.1109/A-SSCC60305.2024.10848802.
- Jacob Ziv and Abraham Lempel. A universal algorithm for sequential data compression. *IEEE Transactions on Information Theory*, 23(3):337–343, 1977. doi: 10.1109/TIT.1977.1055714.
- zlib-ng contributors. zlib-ng: zlib replacement with optimizations for next generation systems, 2025. URL <https://github.com/zlib-ng/zlib-ng>. GitHub repository, v2.3.1.